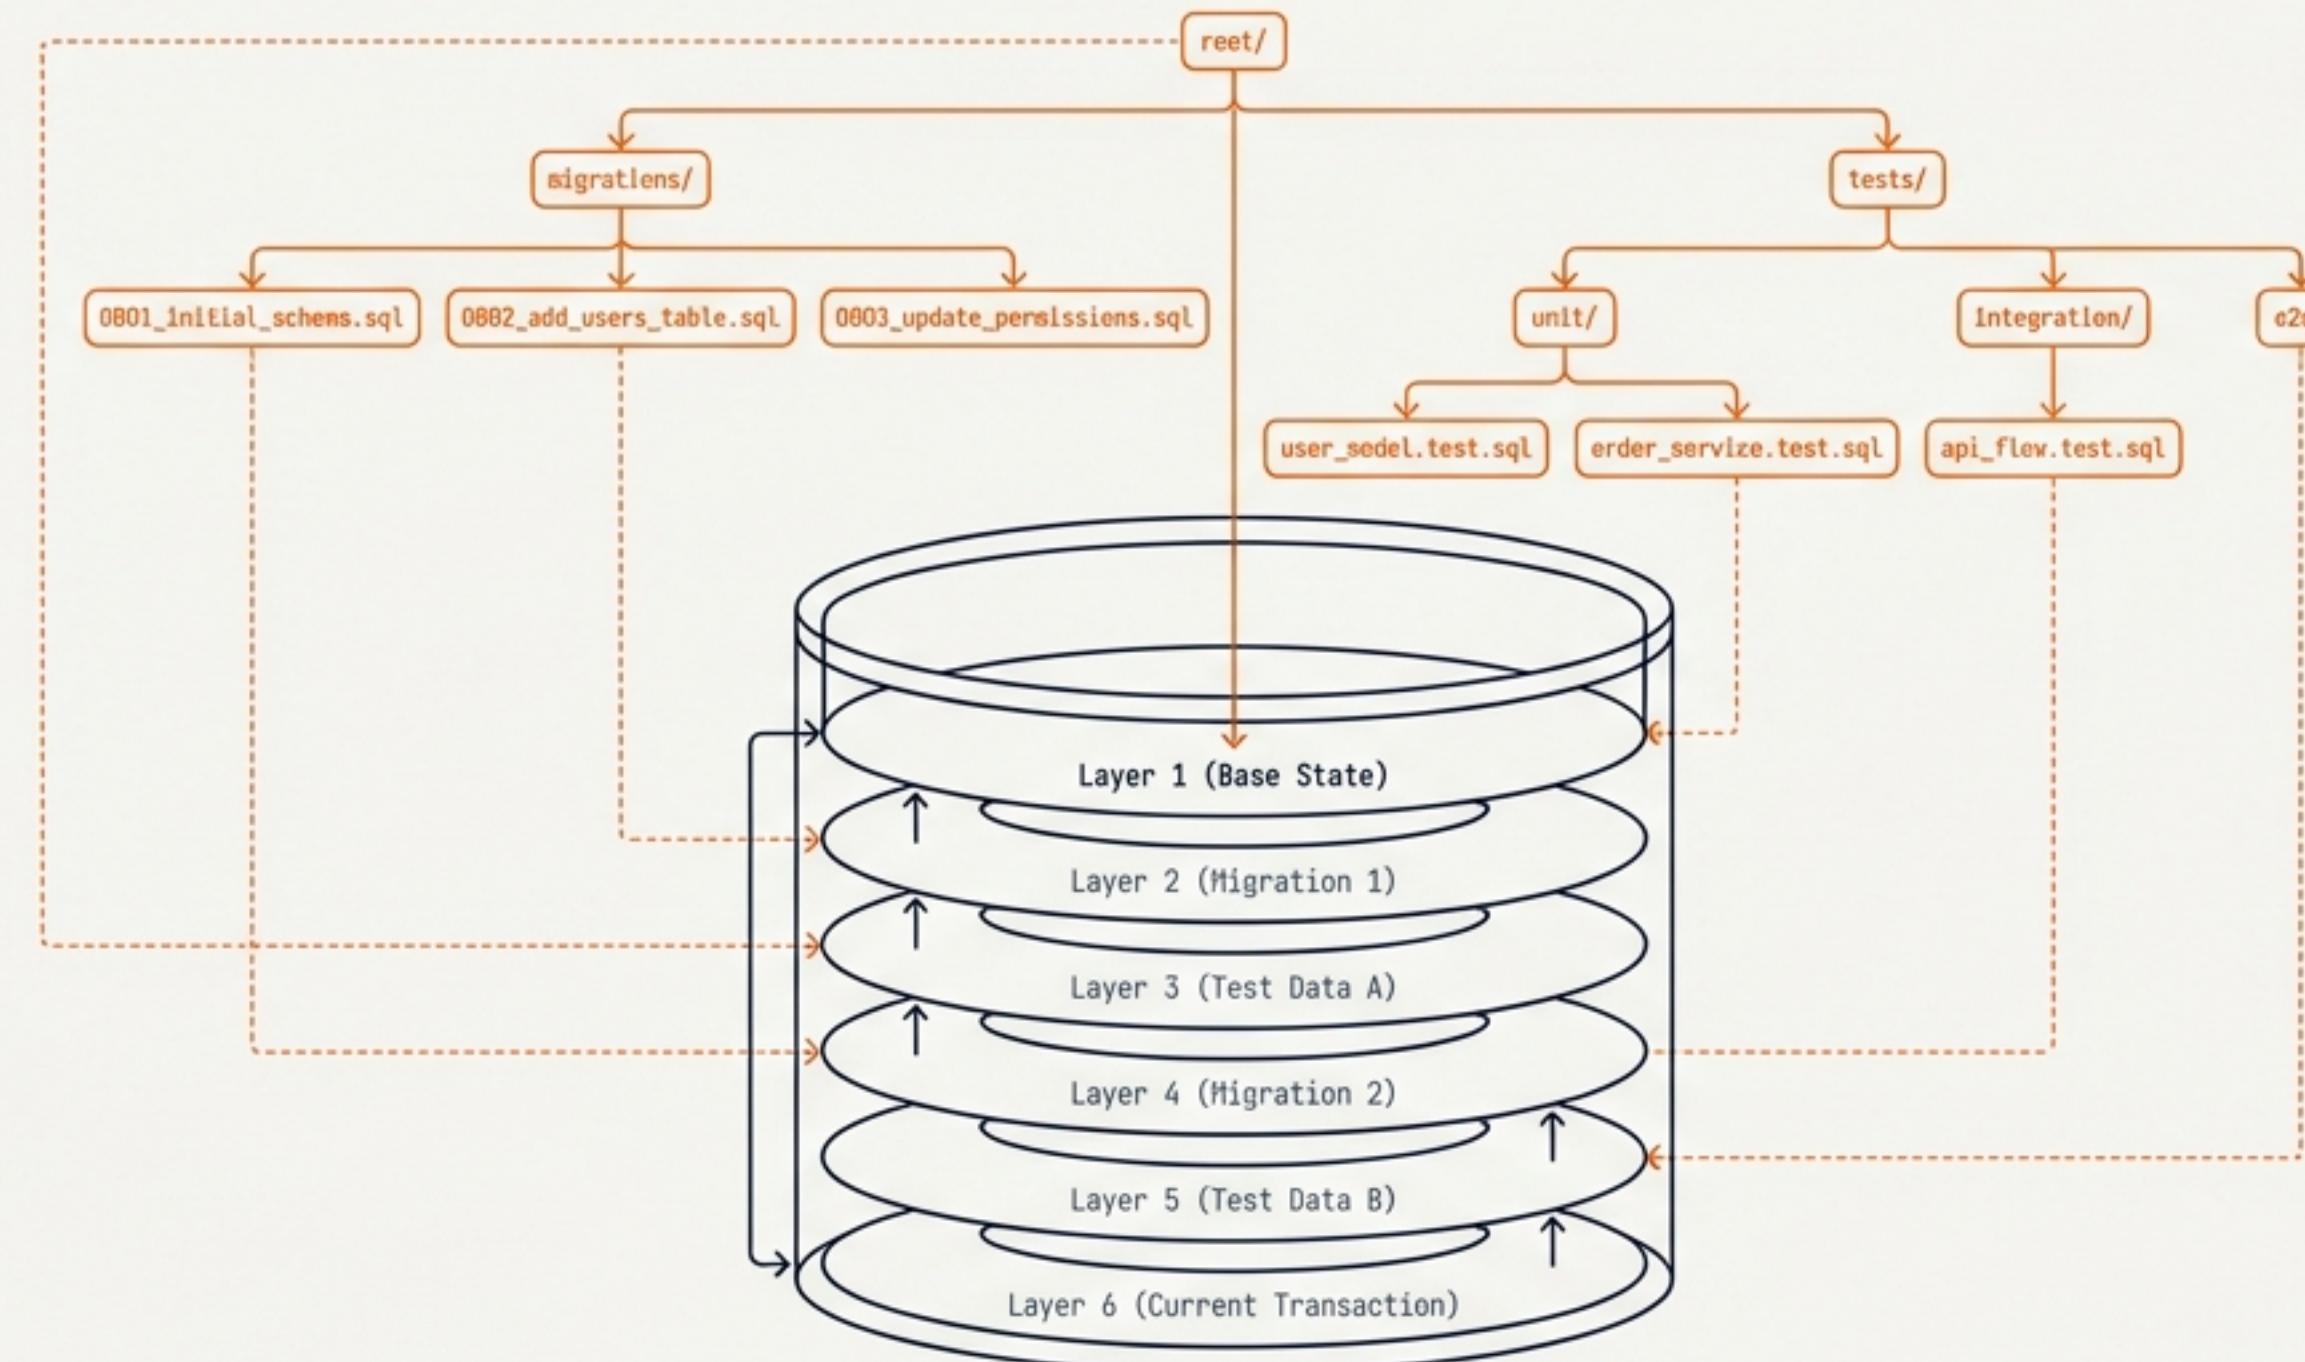


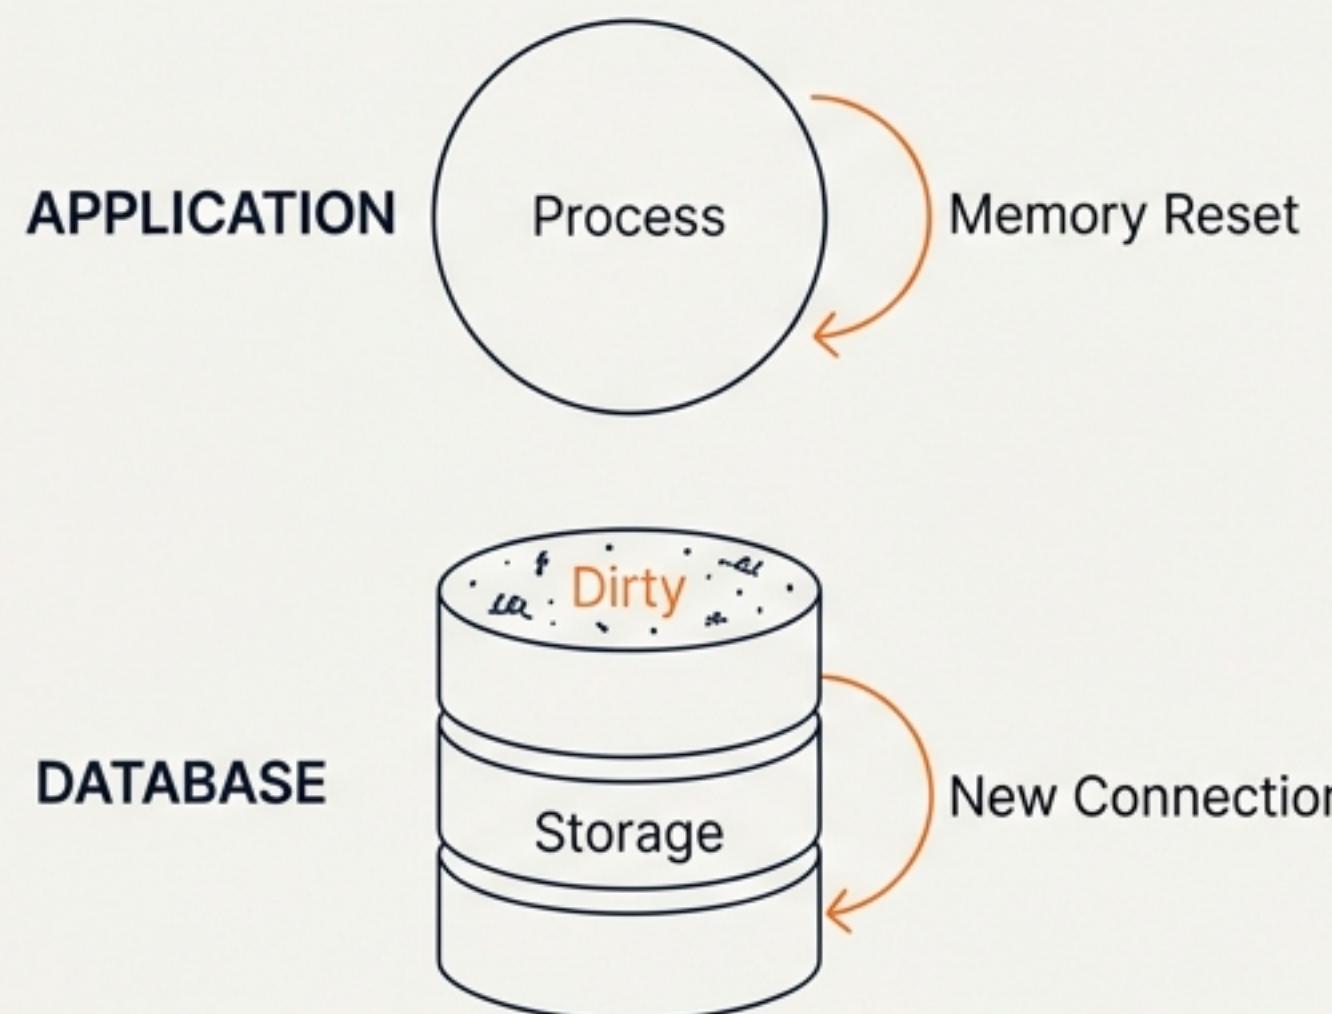
# pgmi: Transactional Testing & Deployment Architecture

A first-principles analysis of zero-side-effect database testing using PostgreSQL primitives



# THE PERSISTENCE PROBLEM: WHY DATABASE TESTING IS HISTORICALLY PAINFUL

## THE CORE TENSION



Applications reset state on restart. Databases persist state by design. This persistence causes "residue" that contaminates tests.

## THE CURRENT LANDSCAPE

<b>Containers (Docker)</b>	Pros Perfect Isolation	Cons High Latency (2-5s startup). Disconnects testing from deployment.
<b>Truncation / Cleanup</b>	Pros Fast Execution	Fragile Maintenance. Cascading deletes are messy. Requires perfect teardown scripts.
<b>ORM Transactions</b>	Pros Easy Integration	Cons App-layer only. Cannot test DDL (schema changes). Ignores DB-internal logic.

**THE INSIGHT: WHAT IF THE DEPLOYMENT TOOL DIDN'T MANAGE THE TRANSACTION, BUT THE SCRIPT DID?**

# THE PRIMITIVES: WHY THIS ONLY WORKS IN POSTGRESQL

pgmi is an orchestration layer for these native capabilities.



## TRANSACTIONAL DDL

Unlike MySQL or Oracle, PostgreSQL allows `CREATE TABLE` and `ALTER TABLE` to be rolled back.

```
BEGIN;  
DROP TABLE users; -- Oops  
ROLLBACK; -- Table users is restored
```



## SAVEPOINTS

Named markers allow nested rollback scopes within a single transaction without aborting the whole operation.

```
SAVEPOINT test_start;  
INSERT INTO ...;  
ROLLBACK TO SAVEPOINT test_start;
```



## SESSION-SCOPED TABLES

`pg\_temp` schema visibility is restricted to the current connection. Data vanishes on disconnect.

```
CREATE TEMP TABLE pgmi_source (...)  
ON COMMIT DROP;
```

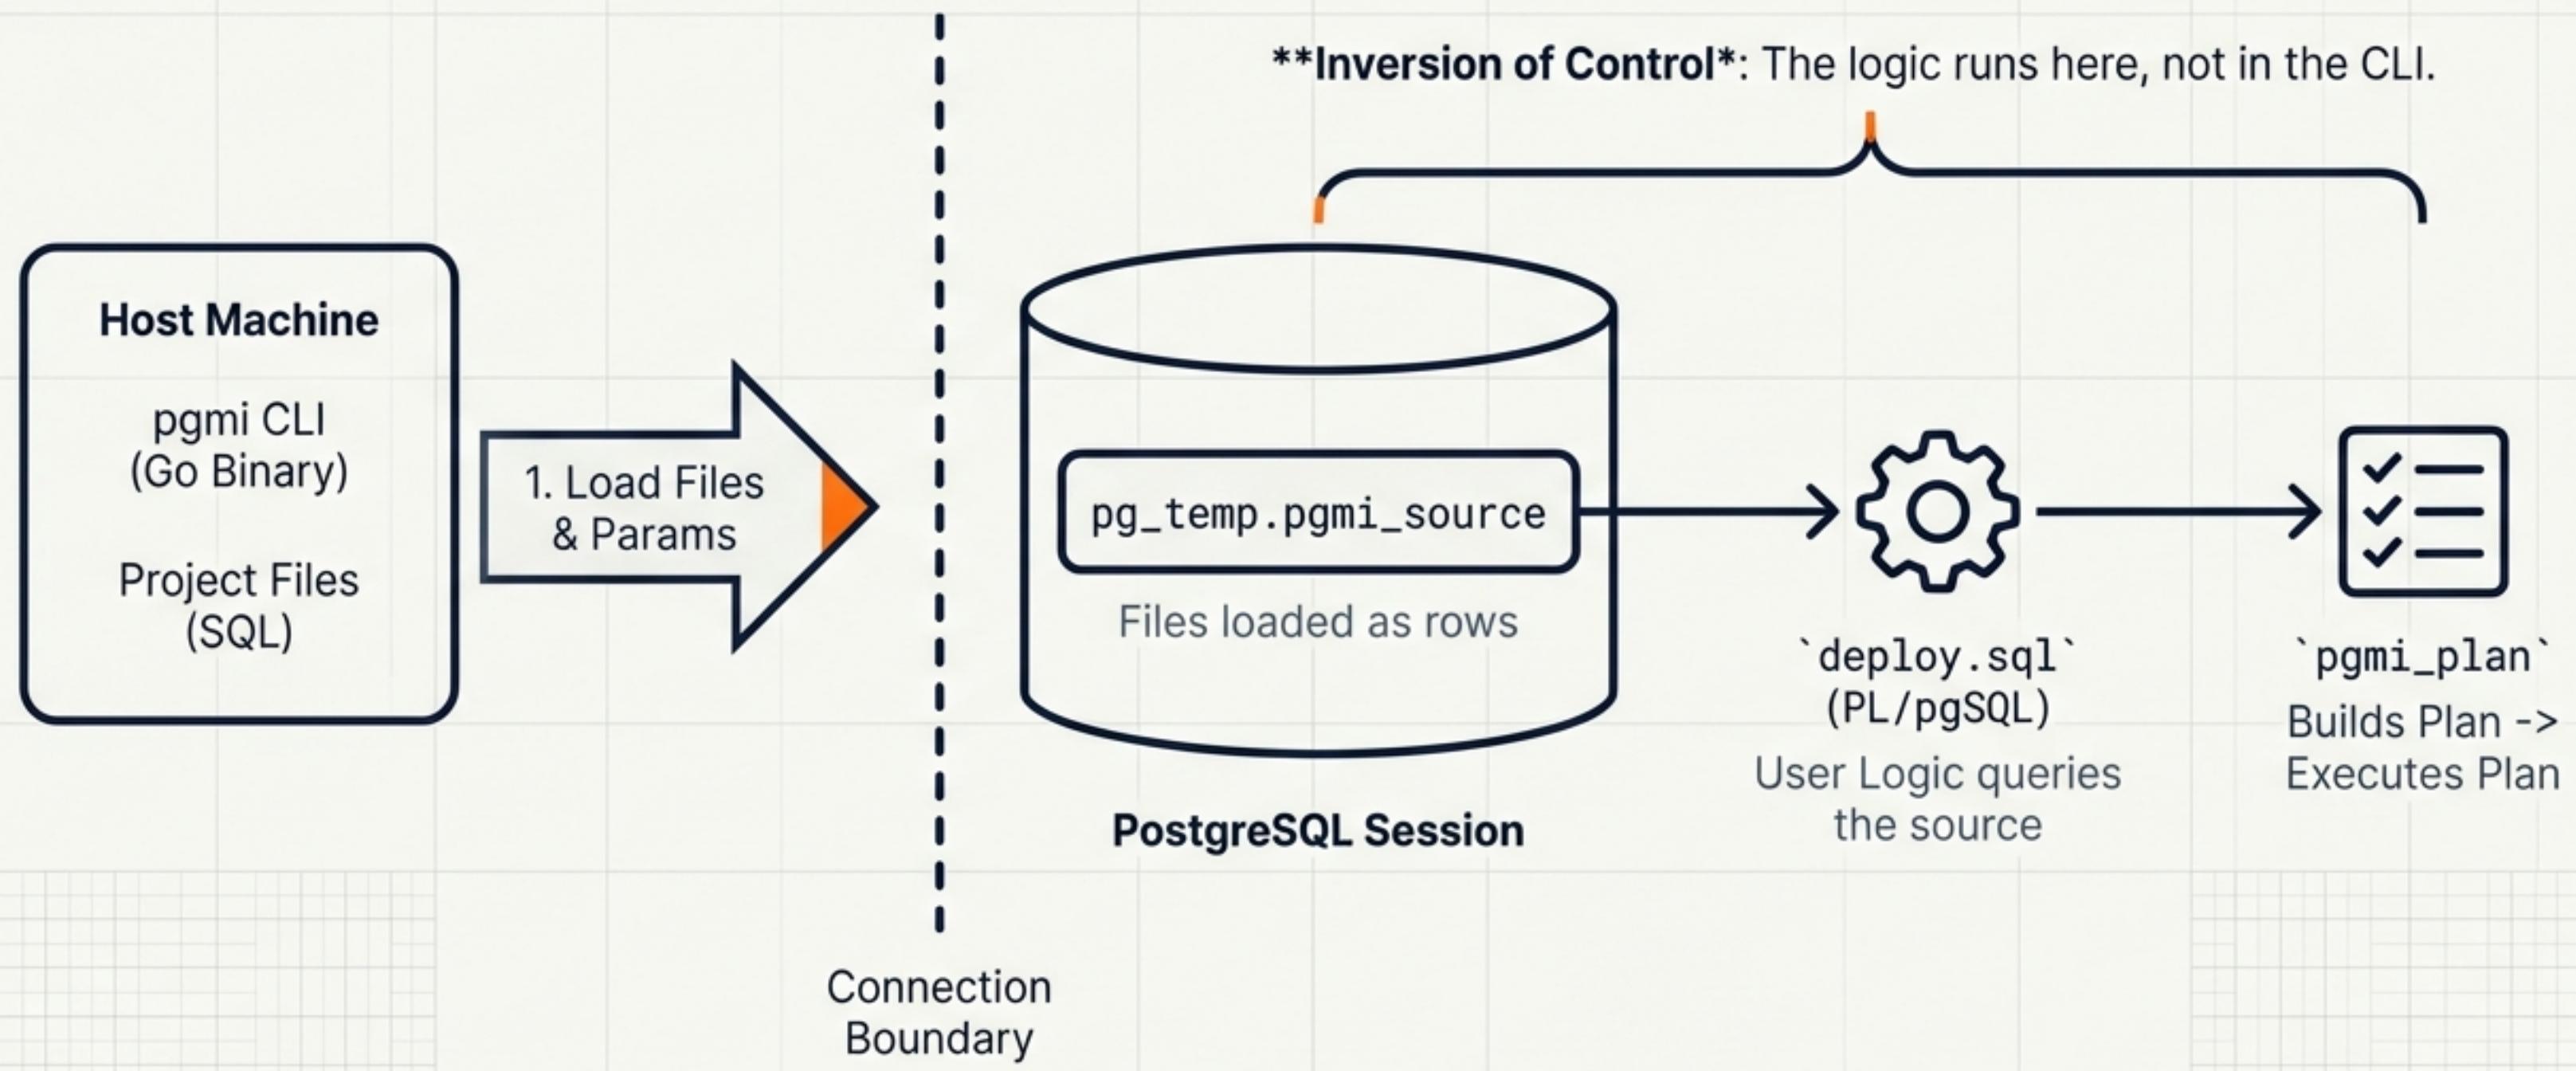


## PL/pgSQL

Turing-complete procedural language allows loops, conditionals, and exception handling inside the DB.

```
DO $$  
BEGIN  
    IF exists THEN ... END IF;  
END $$;
```

# Architecture: PostgreSQL as the Deployment Engine



# The Session API: The Plan is Data

Deployment state is exposed as standard SQL tables in `pg\_temp`.

pgmi_source
path (PK) : text
content : text
checksum : char(64)
directory : text
depth : int

Contains the raw project files loaded from disk.

pgmi_parameter
key (PK) : text
value : text
type : text
is_secret : boolean

Strongly typed CLI arguments injected into the session.

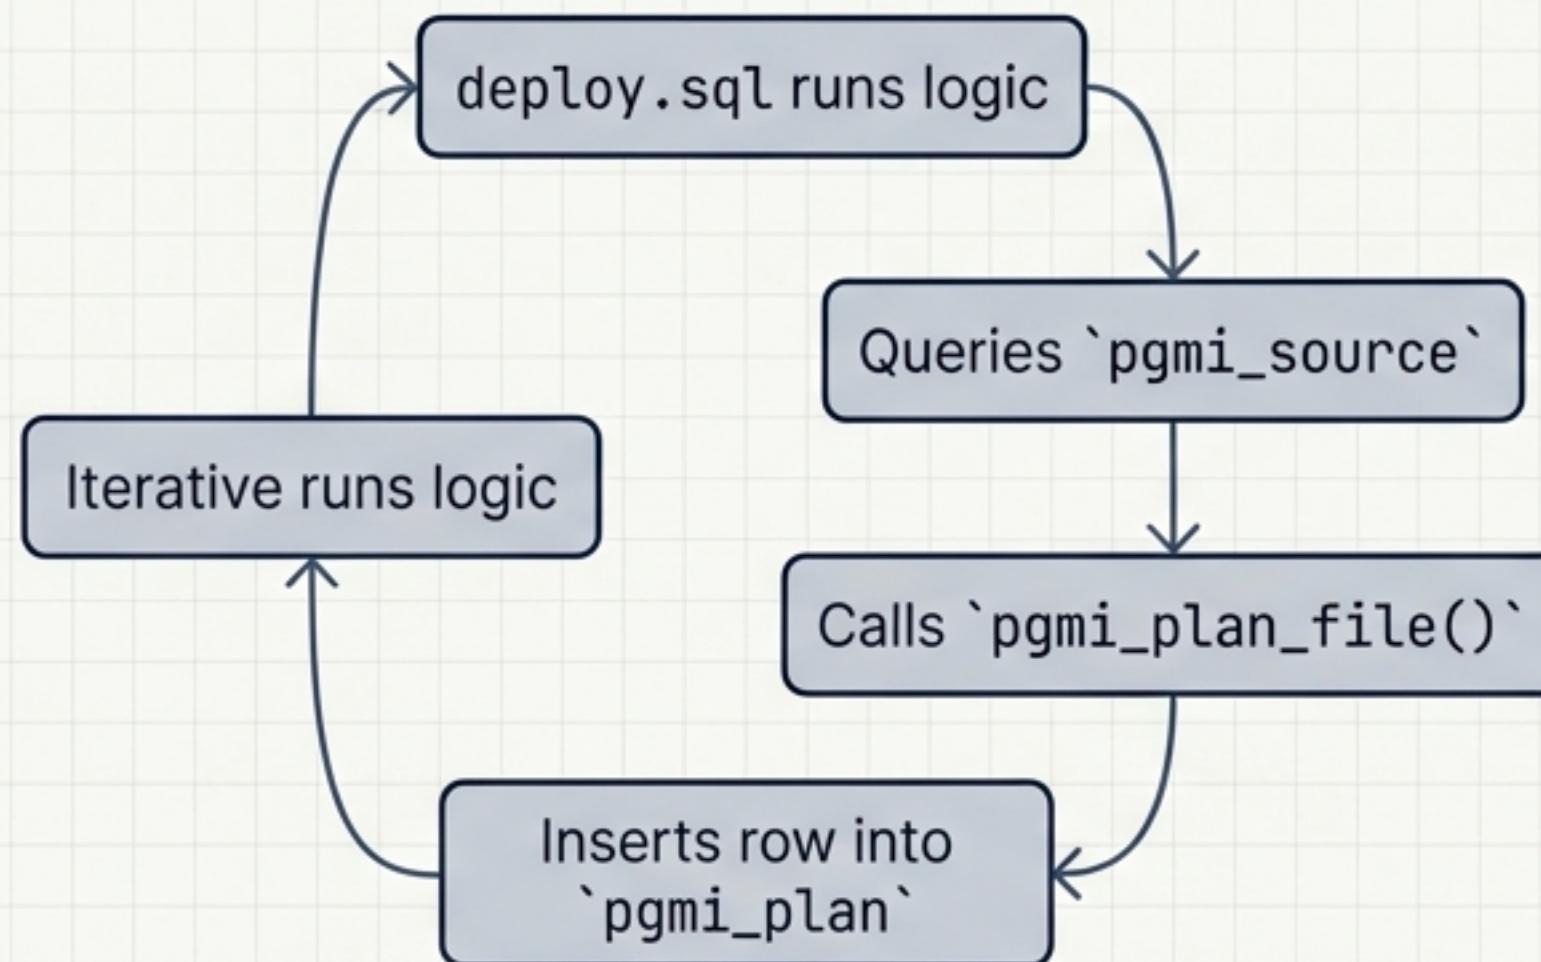
pgmi_plan
ordinal (PK) : serial
command_sql : text

The linear queue of SQL commands to be executed.

```
-- Example: Dynamic Plan Construction
INSERT INTO pgmi_plan (command_sql)
SELECT content FROM pgmi_source
WHERE directory = './migrations/'
ORDER BY name;
```

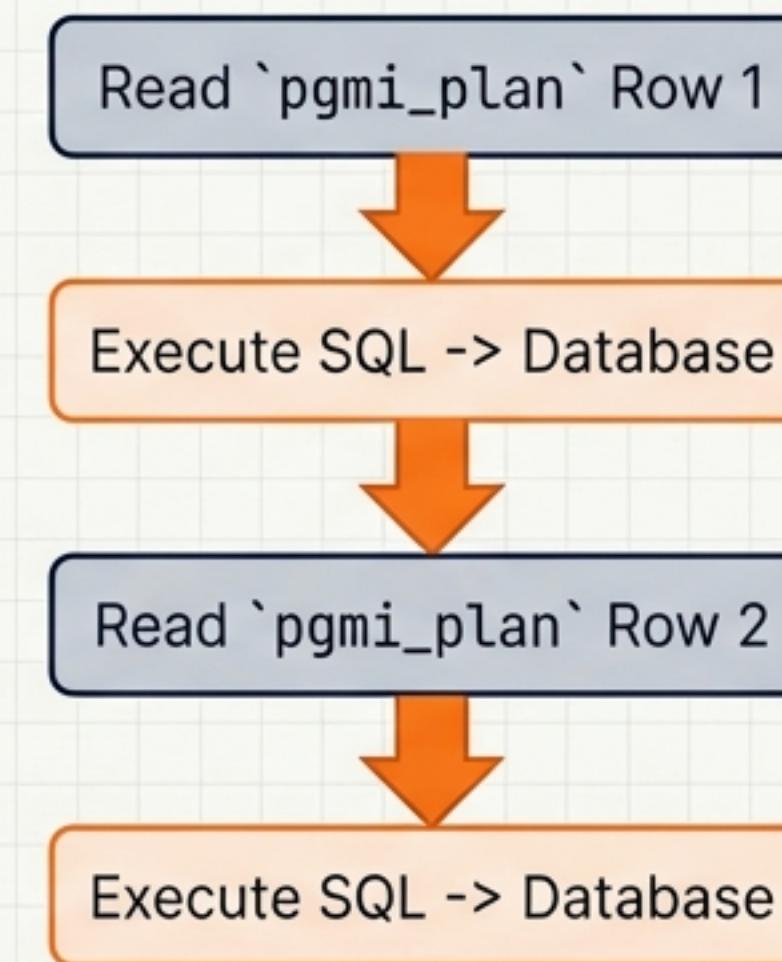
# The Two-Phase Execution Model

## Phase 1: Planning (PL/pgSQL)



No database changes yet. Just data entry.

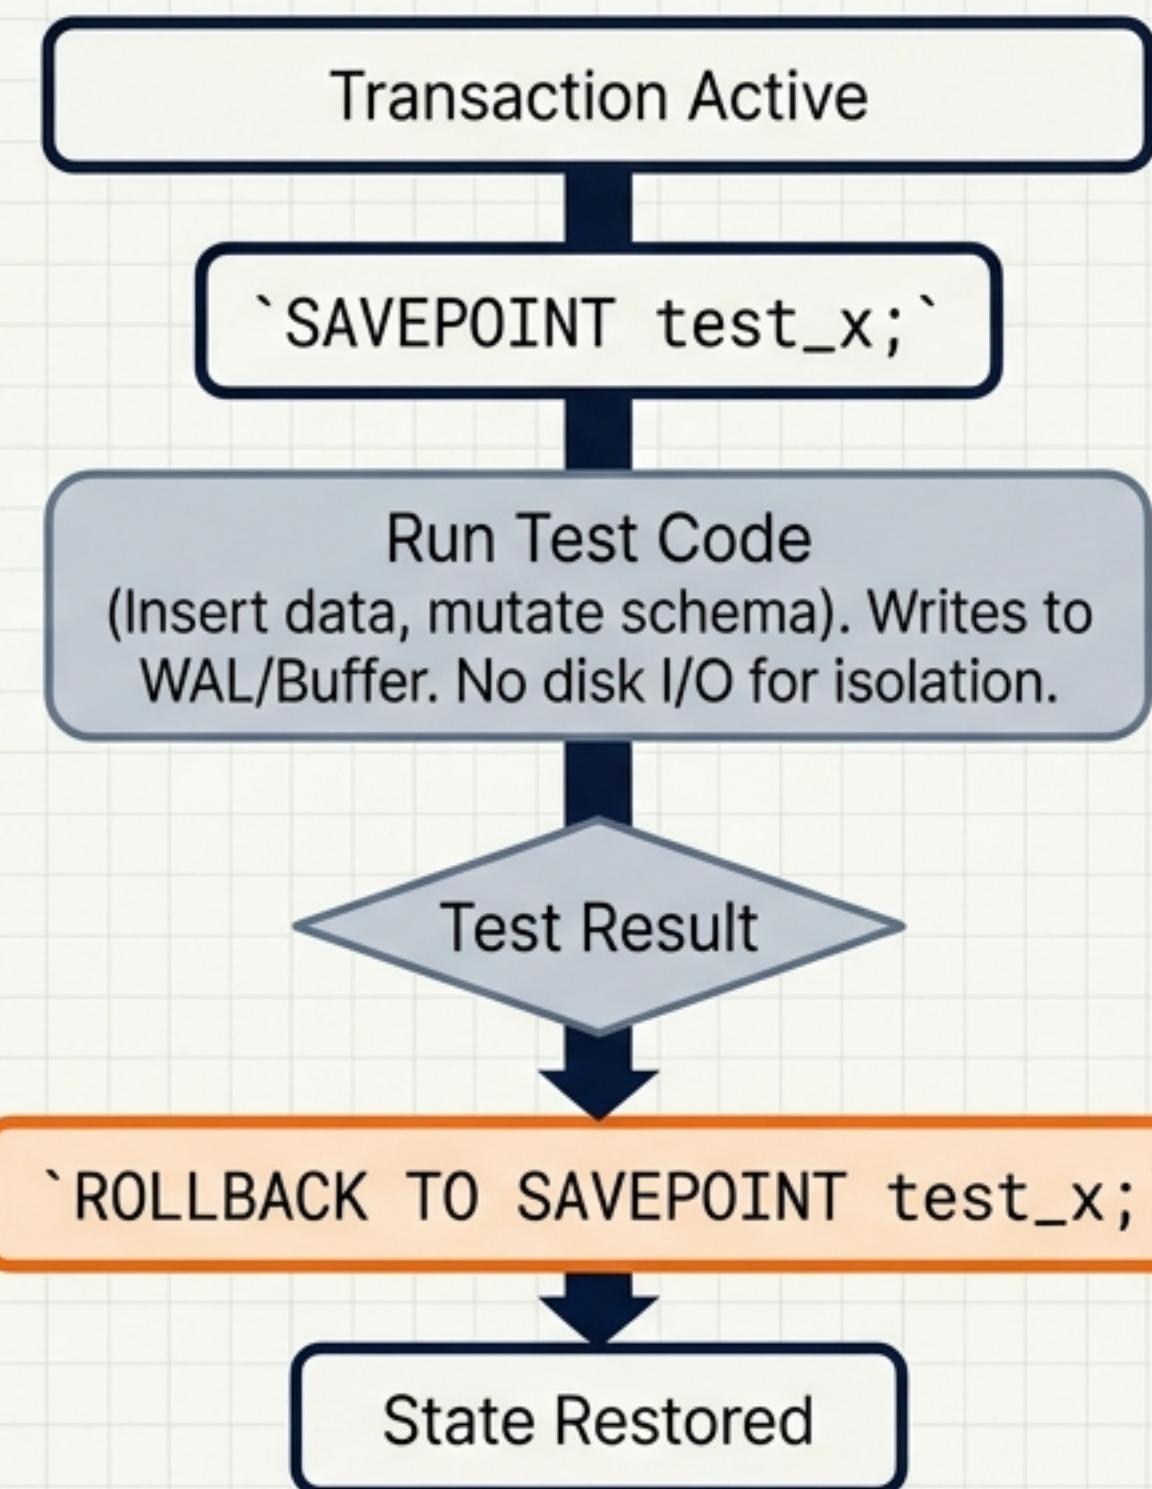
## Phase 2: Execution (Go Binary)



Dumb execution. Runs strings from memory.

This separation enables dynamic test suite construction based on runtime conditions.

# Isolation Strategy: The Savepoint Mechanism



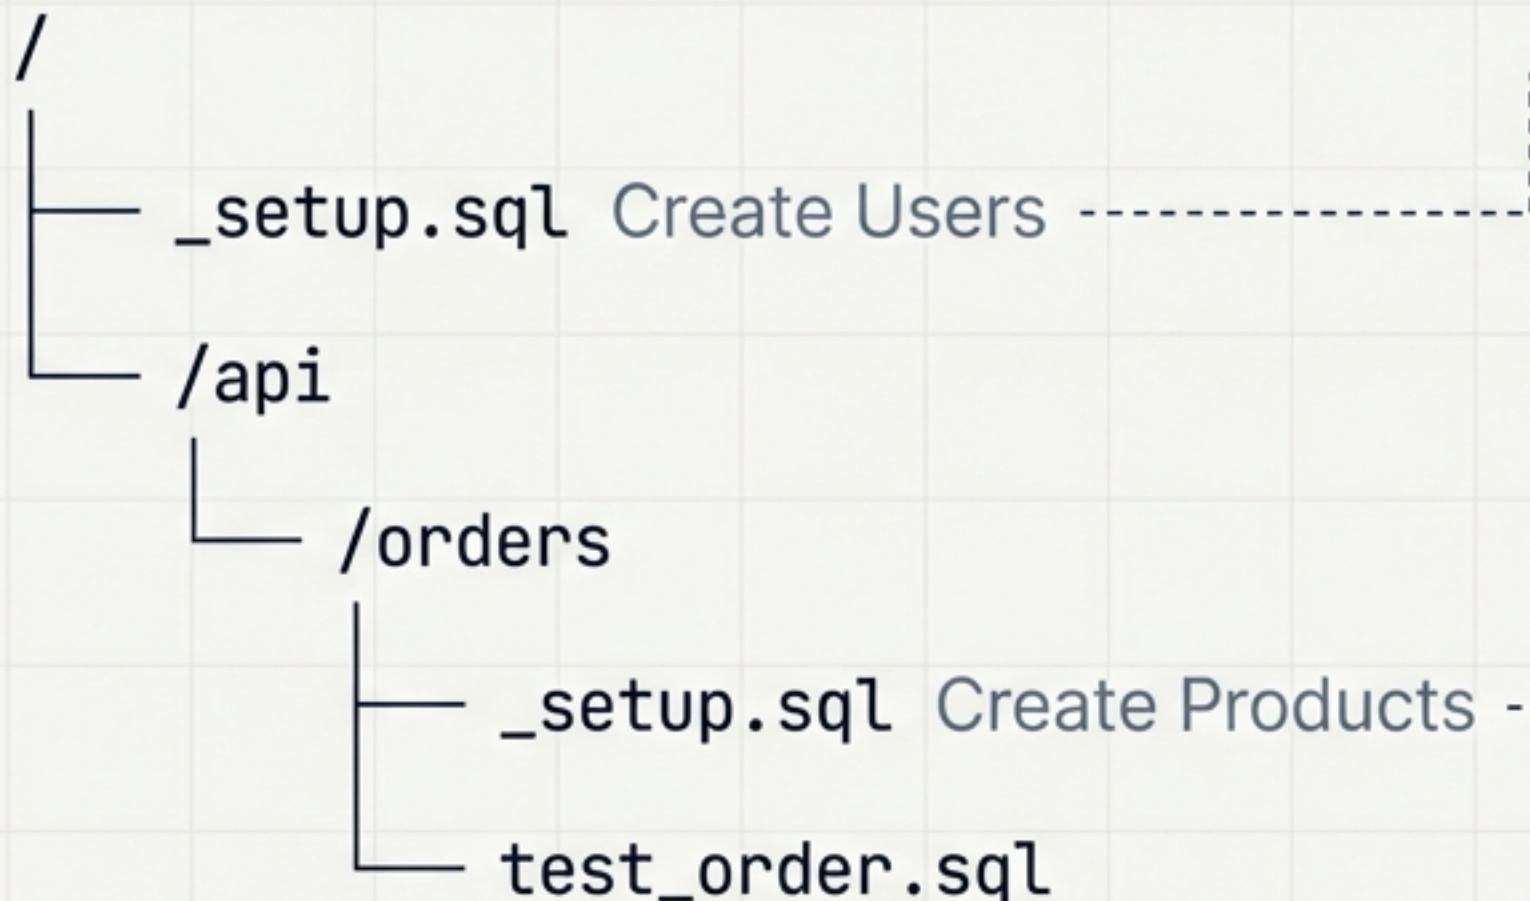
## Performance & Constraints

- **Zero Latency:** No container startup or connection overhead.
- **Metadata Operation:** Savepoints manipulate transaction state, not data pages.
- **The 64-Limit:** PostgreSQL caches 64 subtransactions. `pgmi` avoids overflow by releasing/rolling back *\*before\** starting the next test.  
Active count = Depth, not Total.

# Hierarchical Fixture Composition

The directory tree **\*is\*** the data dependency tree.

## File System (Tree)



## Database Scope (Nested Containers)

SAVEPOINT sp\_root

Users Table Exists

SAVEPOINT sp\_api

(Inherits Users)

SAVEPOINT sp\_orders

Products Table Exists

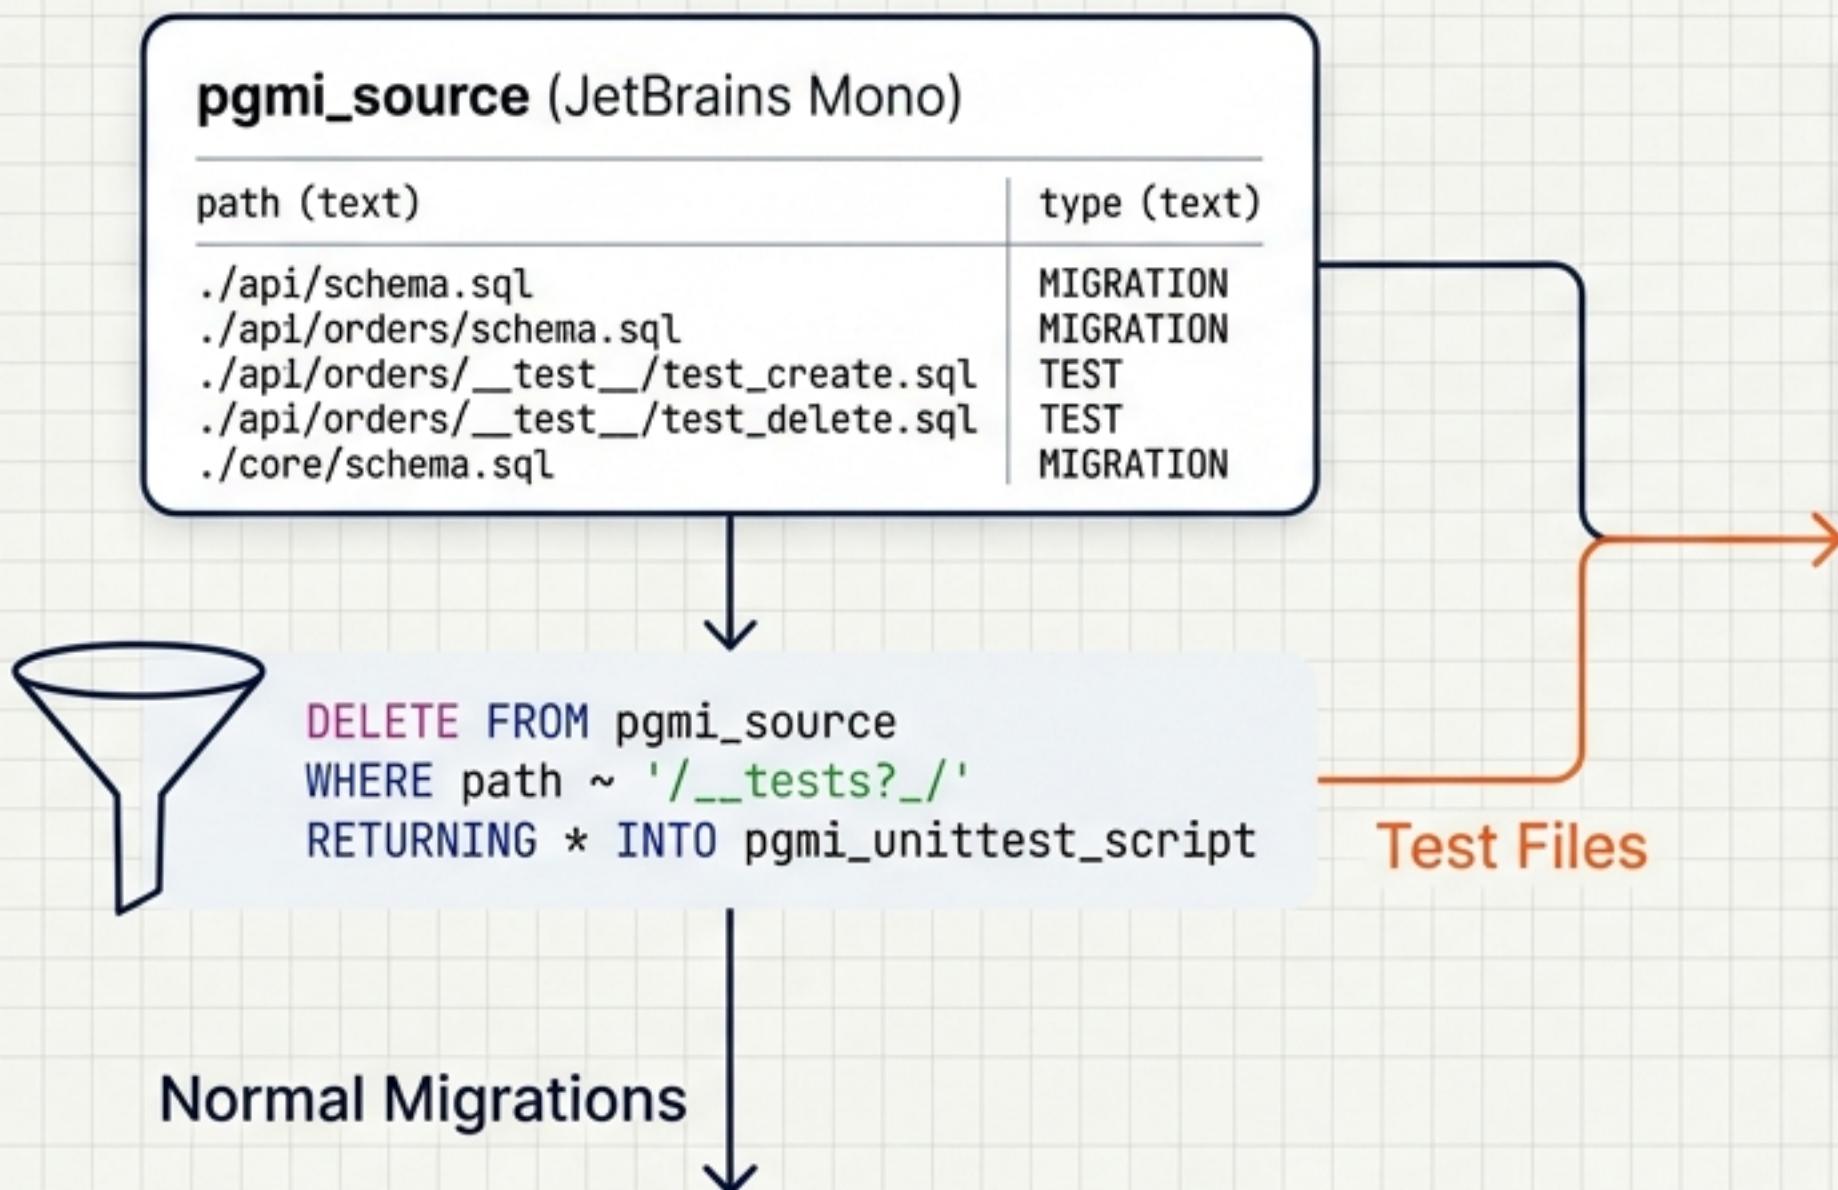
Run `test\_order.sql`

Tests automatically inherit the context of all ancestor setup scripts. Cleanup happens in reverse.

# The Algorithm I: Discovery & Aggregation

Transforms a flat list of files into a navigable tree structure with metadata.

## Step 1: Filtering



## Step 2: Aggregation

**pgmi\_unittest\_vw\_directory**

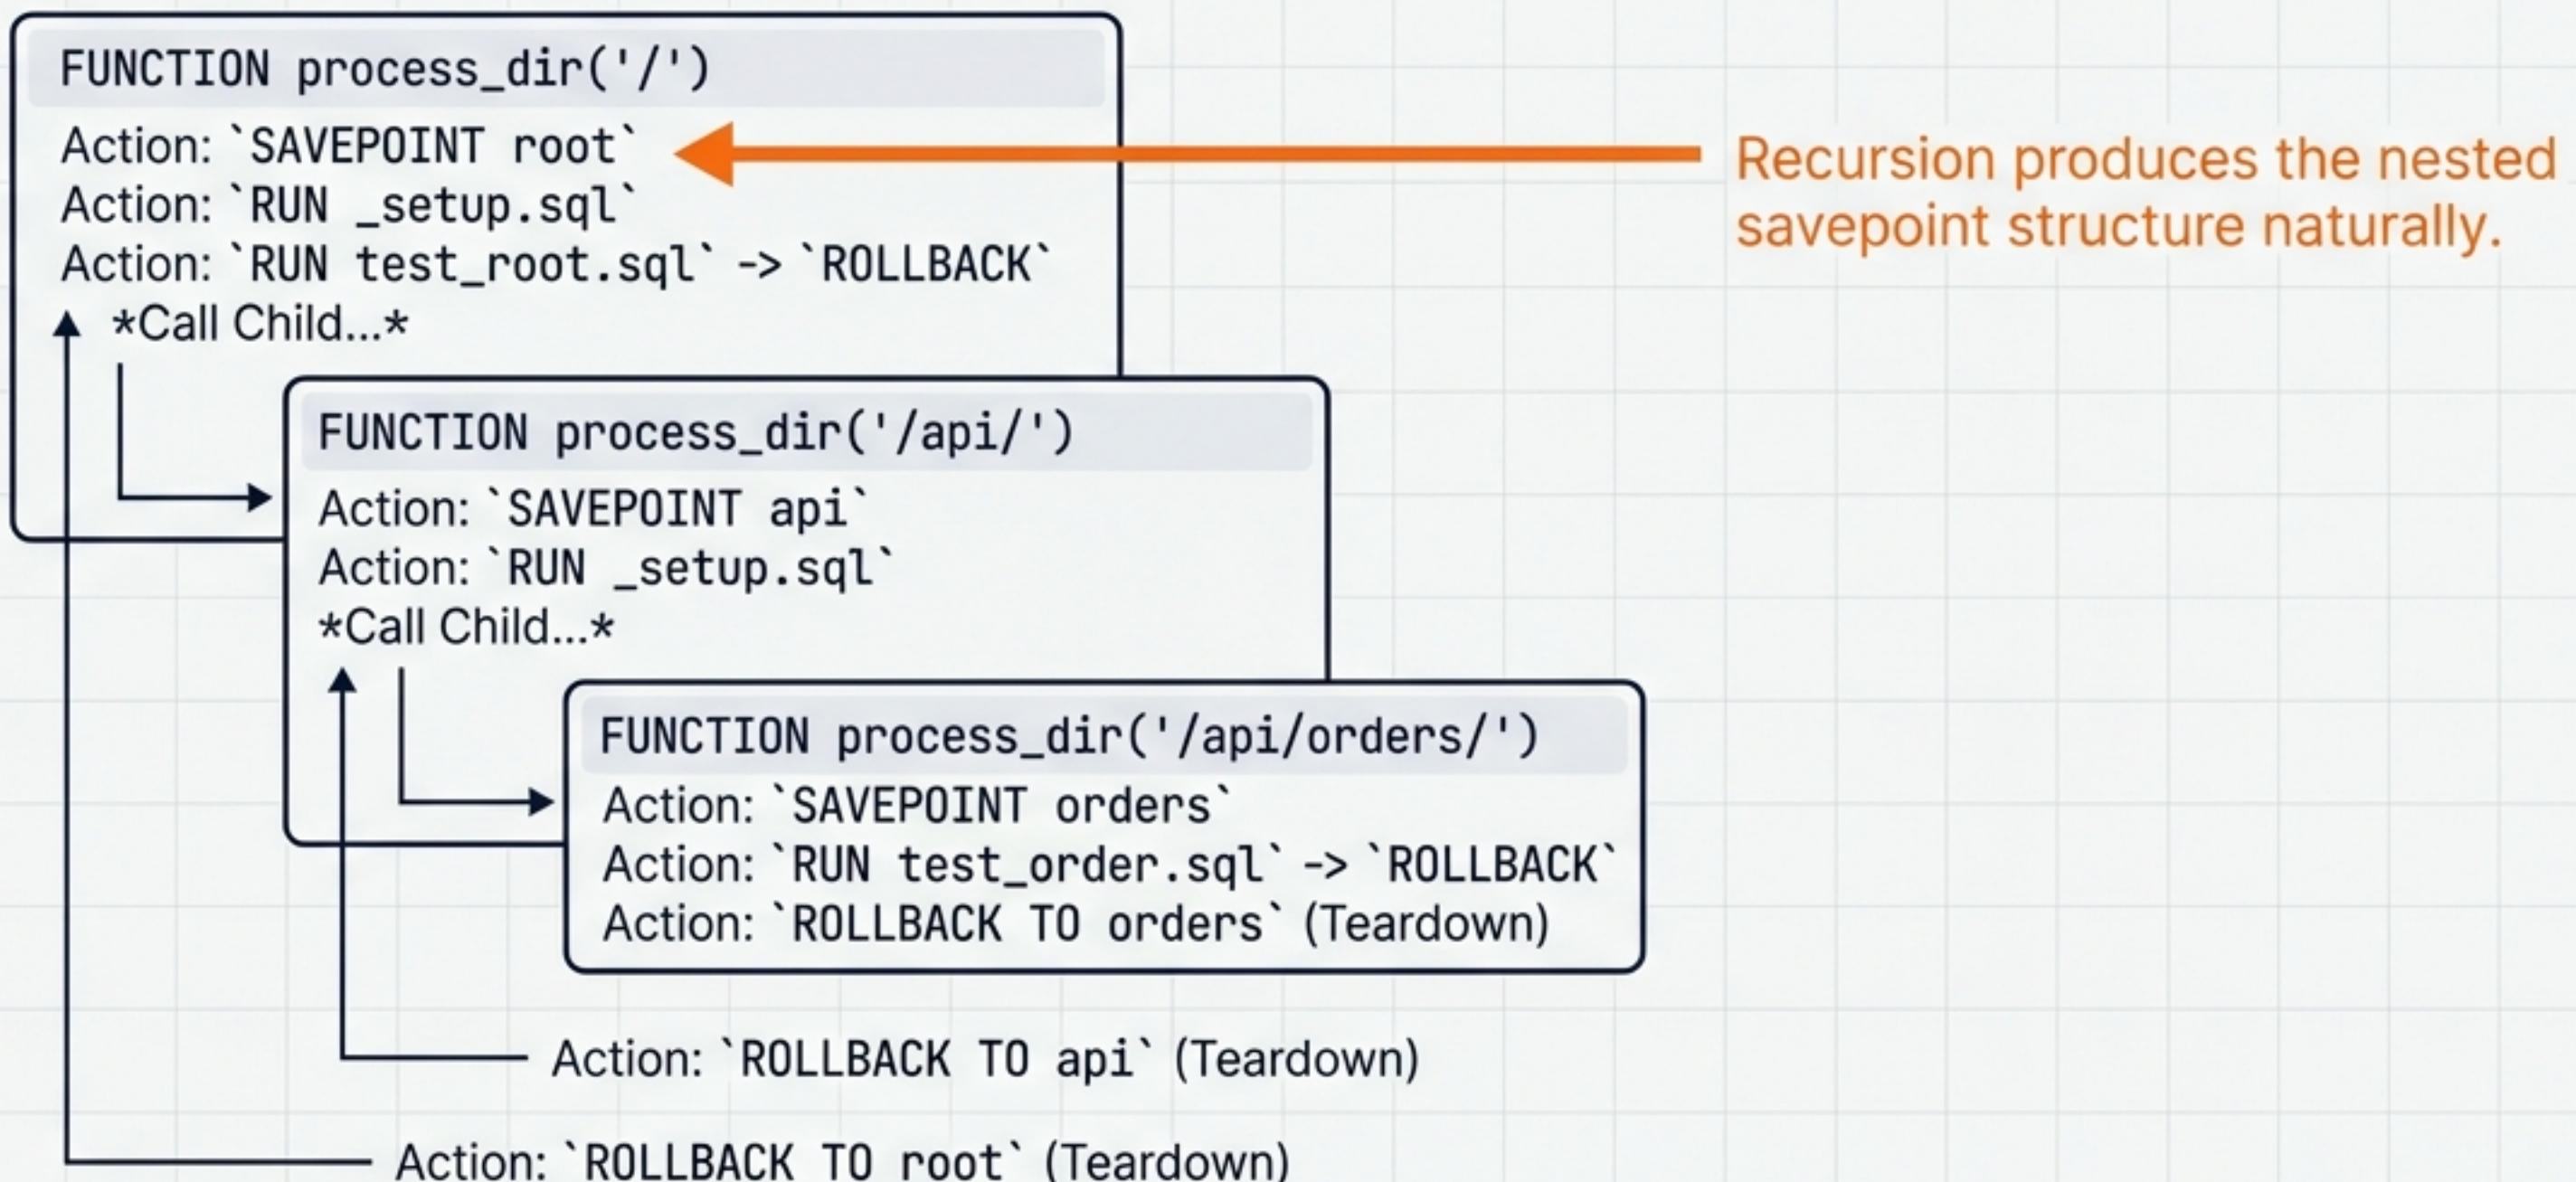
```
{  
  "directory": "./api/orders/",  
  "parent": "./api/",  
  "has_setup": TRUE,  
  "has_teardown": FALSE,  
  "test_scripts": [  
    "test_create.sql",  
    "test_delete.sql"  
  ]  
}
```

Transforms a flat list of files into a navigable tree structure with metadata.

# The Algorithm II: Recursive Traversal

Depth-First Search logic in `pgmi\_unittest\_pvw\_script()`

Call Stack



# The Algorithm III: Plan Materialization

## The Stream

The recursive function yields rows.

Yield  
↓

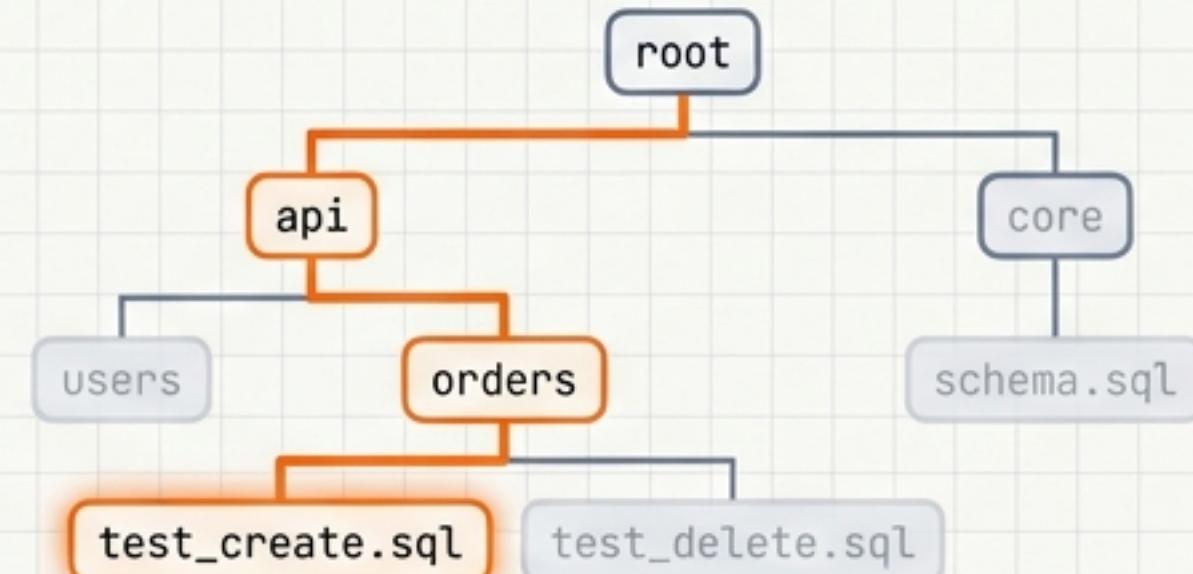
pgmi_plan		
ordinal	command_sql	executable_sql
10	SAVEPOINT sp_001; -- Setup	<pre>CREATE TABLE users (     id SERIAL PRIMARY KEY,     username VARCHAR(50) UNIQUE NOT NULL,     email VARCHAR(100) UNIQUE NOT NULL,     created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP );</pre>

## The Embedding

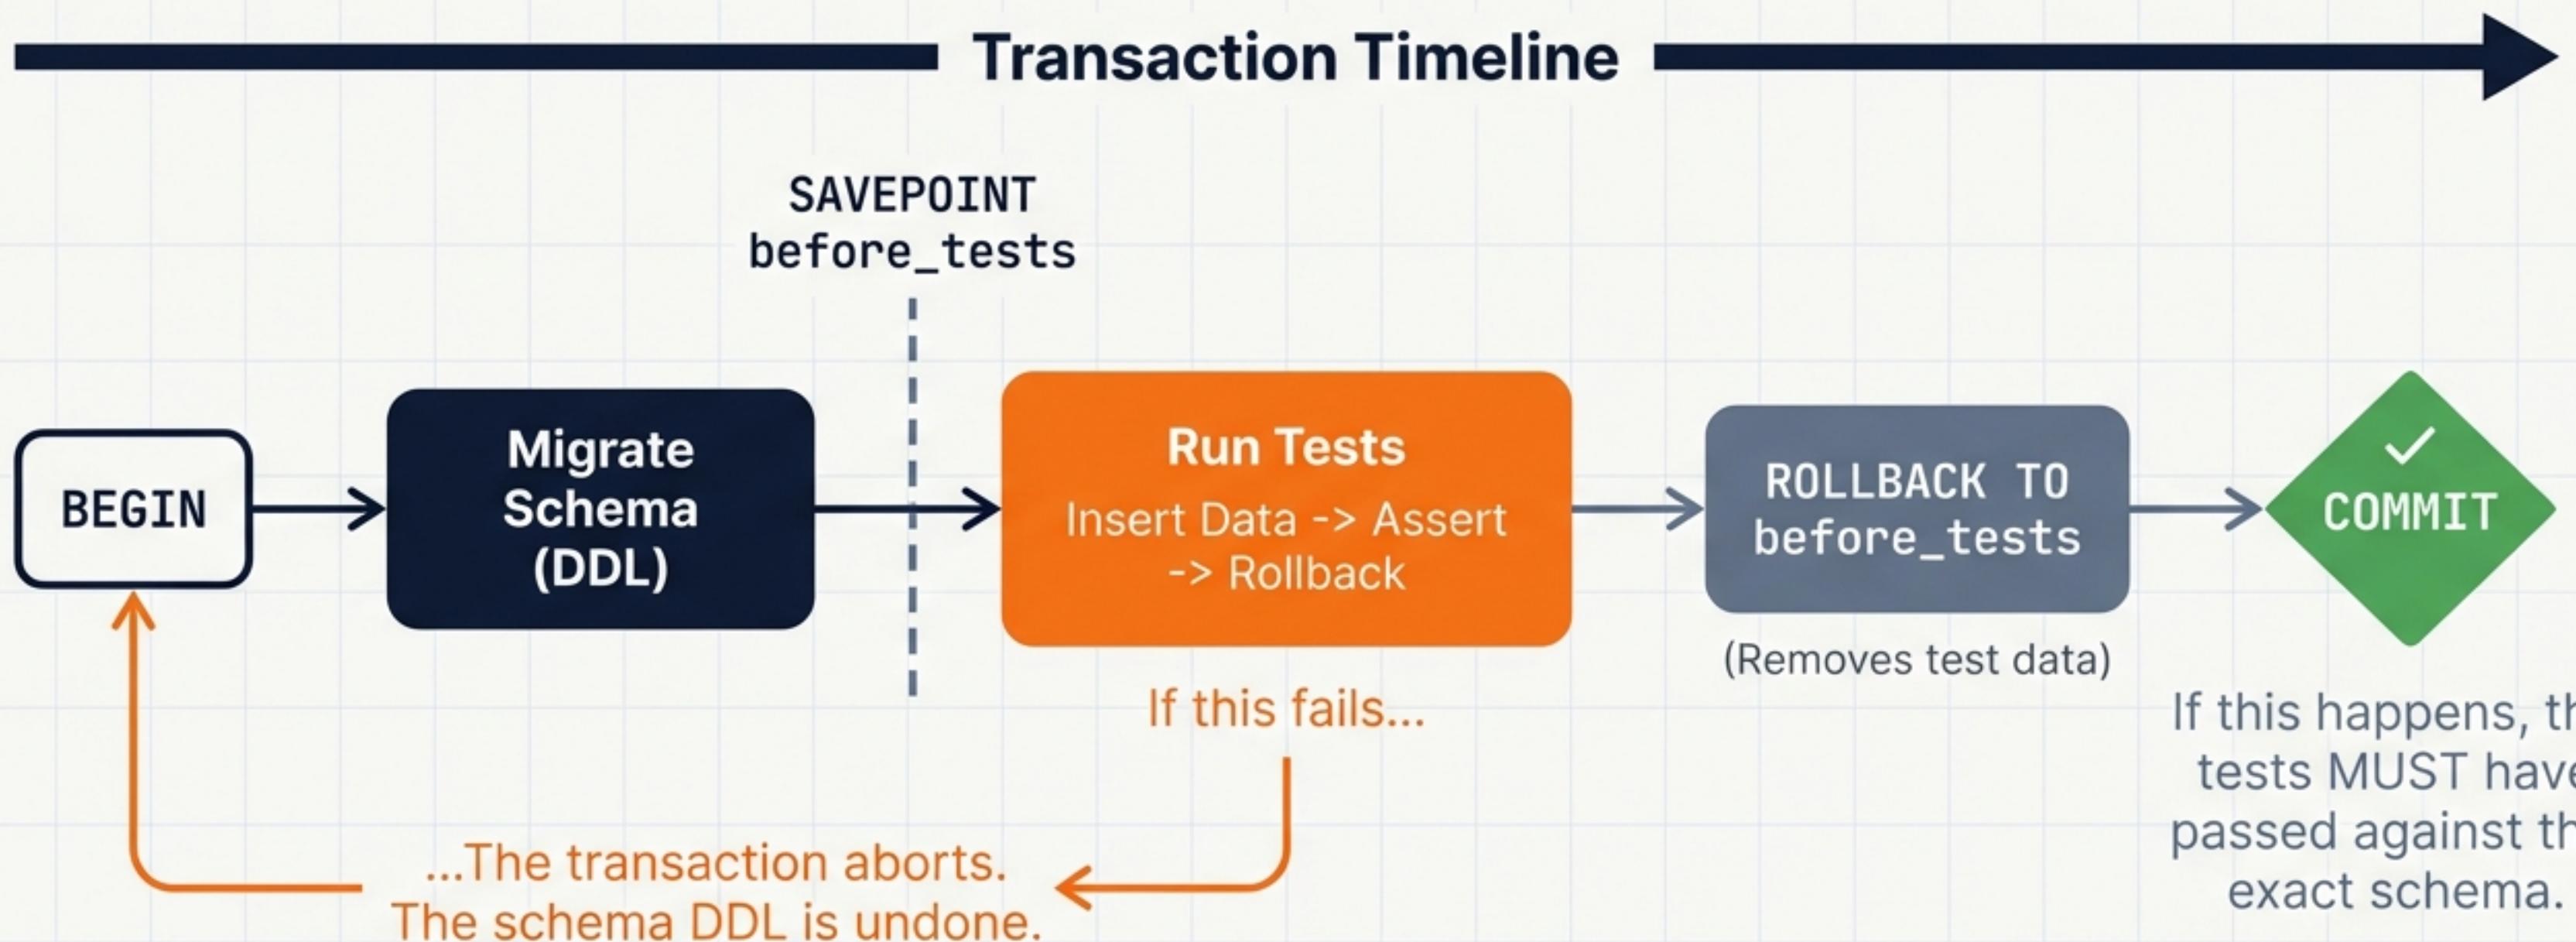
`pgmi` embeds the actual SQL content into the plan.

## Intelligent Filtering

-- When filtering for "/orders/", we preserve parents:  
`WHERE path LIKE parent_path || '%'`



# The Payoff: The 'Gated Deployment' Pattern



**Atomic Deploy-and-Test. Zero bad schemas in production.**

# Compliance & Safety Implications

Meeting PCI-DSS, SOX, and HIPAA requirements with cryptographic certainty.

## Traditional Risk



Deployment without verification.

**Outcome: \$440M Loss.**  
Untested code activated in production.

## Gated Safety

### The Audit Trail

```
> pgmi deploy
[INFO] Migrating schema... OK
[INFO] Running tests...
    - test_orders.sql ... PASS
    - test_users.sql ... PASS
[INFO] Tests passed. Committing transaction.
Exit Code: 0
```

- Evidence of testing *during* change event.
- Atomic guarantee.
- Audit trail via `RAISE NOTICE`.

**The exit code is the compliance artifact.**

# Tradeoffs & Constraints

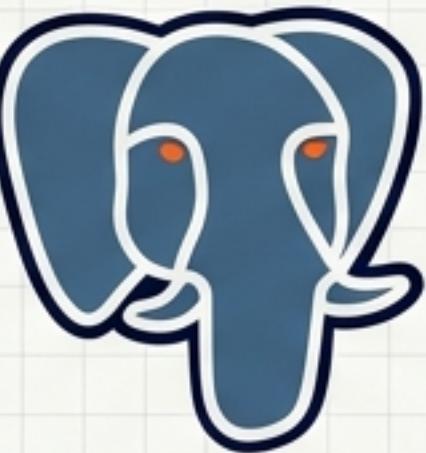
Where pgmi hits the boundaries of the model.



## No Parallel Execution.

Reason: pg\_temp tables and Transactions are bound to a single session.

**Impact:** Test suite scaling is linear, not horizontal.



## PostgreSQL Only

Reason: Relies on Transactional DDL and Savepoints behavior specific to Postgres.



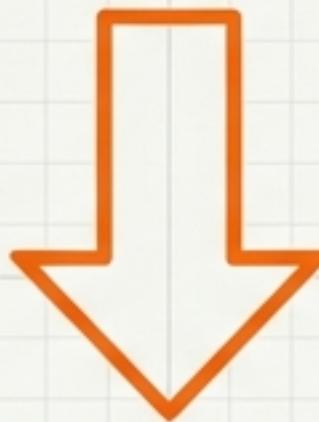
## No Side-Effects.

Reason: Cannot rollback external API calls (HTTP, S3) via ROLLBACK.

# Summary: Evolutionary Database Design

pgmi moves deployment logic from external tools (Java/Go/Perl) into the database itself (PL/pgSQL).

**Test-Then-Deploy**  
(Race conditions, drift)



**Deploy-And-Test  
(Atomic guarantee)**

“It requires you to think about transaction boundaries explicitly, but gives you the power to control them completely.”

Open Source: [github.com/vvka-141/pgmi](https://github.com/vvka-141/pgmi) | Documentation: [pgmi.dev](https://pgmi.dev)